

Analyzing MLIR GPU Dialect for Neural Network Workloads and Proposing Optimizations

KV Tuhil, Aditya Garg, Anirudh Saikrishnan

November 2025

1 Introduction

Deep Neural Networks (DNNs) like ResNet-50 rely on GPUs for intensive tensor computations, a process increasingly managed by the LLVM-based MLIR framework and its hardware-aware GPU dialect. While MLIR effectively handles abstractions such as thread mapping and memory hierarchy to generate CUDA or ROCm kernels, there remains significant potential to improve how these workloads are compiled. This project focuses on analyzing the MLIR GPU dialect generated for ResNet-50 to identify specific performance bottlenecks such as inefficiencies in synchronization or memory usage and proposes novel optimization strategies to enhance execution efficiency.

2 Overview

The project workflow proceeded through the following stages:

1. **Investigation of MLIR Dialects:** We investigated various MLIR dialects, specifically focusing on the SCF (Structured Control Flow) parallel and GPU dialects. The primary objective was to identify opportunities to improve thread block mapping, memory coalescing, and tiling strategies.
2. **Analysis of Lowering Pipeline:** We analyzed the dialect lowering process across various passes, utilizing a convolution kernel as a representative example to trace the transformation of the intermediate representation.
3. **Modification of Standard Passes:** We modified the `ParallelLoopPass` to implement an optimized thread mapping policy, aimed at increasing GPU occupancy.
4. **Integration with IREE:** We transitioned to the IREE compiler infrastructure to validate the modified pass. However, we discovered that IREE utilizes its own custom lowering passes and effectively bypasses the standard SCF parallel dialect.
5. **Optimization and Results:** Consequently, we shifted focus to the IREE LLVMGPU kernel configuration. By optimizing tile sizes and thread mapping parameters, we successfully reduced the kernel runtime from 19.7 ms to 6.84 ms.

3 MLIR Dialect Walkthrough

To analyze the lowering pipeline from high-level IR to the GPU dialect and finally to LLVM, we utilized a convolution kernel as a representative benchmark.

- **Linalg Dialect:** The process begins with a convolution operation in the `linalg` dialect. This operation accepts a $512 \times 16 \times 16$ input matrix and a $512 \times 3 \times 3$ kernel, producing a $512 \times 7 \times 7$ output matrix.

```
%436 = linalg.conv_2d_nchw_fchw [
    dilations = dense<1> : vector<2xi64>, strides = dense<2> : vector<2xi64>
]
ins(%padded_354, %cst_136 : tensor<1x512x16x16xf32>,
tensor<512x512x3x3xf32>) outs(%435 : tensor<1x512x7x7xf32>) -> tensor<1x512x7x7xf32>
```

Figure 1: Snippet of the Convolution operation in the Linalg dialect.

- **Lowering to SCF Parallel:** Subsequent lowering to the `scf` (Structured Control Flow) parallel dialect results in nested loops where the first three loops are parallelized, while the remainder are sequential. Notably, the step size for the parallel loops is set to 1. This default configuration restricts the thread block size to 1, which severely limits GPU occupancy.

```
scf.parallel (%arg160, %arg161, %arg162) = (%c0, %c0, %c0) to (%c512, %c7, %c7) step (%c1, %c1, %c1) {
    scf.for %arg163 = %c0 to %c512 step %c1 {
        scf.for %arg164 = %c0 to %c3 step %c1 {
            scf.for %arg165 = %c0 to %c3 step %c1 {
                %161 = affine.apply #map(%arg161, %arg164)
                %162 = affine.apply #map(%arg162, %arg165)
                %163 = memref.load %alloc_237[%c0, %arg163, %161, %162] : memref<1x512x16x16xf32>
                %164 = memref.load %alloc_132[%arg160, %arg163, %arg164, %arg165] : memref<512x512x3x3xf32>
                %165 = memref.load %alloc_241[%c0, %arg160, %arg161, %arg162] : memref<1x512x7x7xf32>
                %166 = arith.mulf %163, %164 : f32
                %167 = arith.addf %165, %166 : f32
                memref.store %167, %alloc_241[%c0, %arg160, %arg161, %arg162] : memref<1x512x7x7xf32>
            }
        }
    }
    scf.reduce
}
```

Figure 2: Lowered SCF Parallel loops showing unit step sizes.

- **GPU Kernel Mapping:** This IR is mapped to a GPU kernel launch configuration. As shown in the snippet below, the resulting block size is $\{1, 1, 1\}$. This configuration prevents memory coalescing and fails to utilize the GPU’s parallelism effectively.

```
gpu.module @resnet50_kernel_434 {
    gpu.func @resnet50_kernel(%arg0: index, %arg1: index, %arg2: memref<1x512x16x16xf32>, %arg3: memref<512x512x3x3xf32>,
                           %arg4: memref<1x512x7x7xf32>, %arg5: index, %arg6: index)
        kernel attributes {known_block_size = array<i32: 1, 1, 1>} {
            %block_id_x = gpu.block_id_x
            %block_id_y = gpu.block_id_y
            %block_id_z = gpu.block_id_z
            %0 = affine.apply #map()(%arg0, %arg1, %block_id_x)
            %1 = affine.apply #map()(%arg0, %arg1, %block_id_y)
            %2 = affine.apply #map()(%arg0, %arg1, %block_id_z)
            scf.for %arg7 = %arg1 to %arg6 step %arg0 {
                scf.for %arg8 = %arg1 to %arg5 step %arg0 {
                    scf.for %arg9 = %arg1 to %arg5 step %arg0 {
                        scf.for %arg10 = %arg1 to %arg5 step %arg0 {
                            %3 = affine.apply #map1(%arg8)[%arg0, %arg1, %block_id_y]
                            %4 = affine.apply #map1(%arg9)[%arg0, %arg1, %block_id_z]
                            %5 = memref.load %arg2[%arg1, %arg7, %3, %4] : memref<1x512x16x16xf32>
                            %6 = memref.load %arg3[%0, %arg7, %arg8, %arg9] : memref<512x512x3x3xf32>
                            %7 = memref.load %arg4[%arg1, %0, %1, %2] : memref<1x512x7x7xf32>
                            %8 = arith.mulf %5, %6 : f32
                            %9 = arith.addf %7, %8 : f32
                            memref.store %9, %arg4[%arg1, %0, %1, %2] : memref<1x512x7x7xf32>
                        }
                    }
                }
            }
            gpu.return
        }
    }
```

Figure 3: Generated GPU Kernel with inefficient block size mapping.

- **Optimization Strategy:** To address this bottleneck, we modified the mapping policy within the `ParallelLoopPass`. We implemented a greedy strategy to maximize thread utilization. The approach is defined as follows:

- Tuhil please write it

```
%c0_1711 = arith.constant 0 : index
%c1_1712 = arith.constant 1 : index
%c32_1713 = arith.constant 32 : index
%c8_1714 = arith.constant 8 : index
%1199 = arith.muli %c1, %c1_1712 : index
%1200 = arith.muli %c1, %c32_1713 : index
%1201 = arith.muli %c1, %c8_1714 : index
scf.parallel (%arg160, %arg161, %arg162) = (%c0, %c0, %c0) to (%c512, %c7, %c7) step (%1199, %1200, %1201) {
    %1414 = affine.min #map(%1200, %c7, %arg161)
    %1415 = affine.min #map(%1201, %c7, %arg162)
    scf.parallel (%arg163, %arg164, %arg165) = (%c0_1711, %c0_1711, %c0_1711) to (%1199, %1414, %1415) step (%c1, %c1, %c1) {
        %1416 = arith.addi %arg163, %arg160 : index
        %1417 = arith.addi %arg164, %arg161 : index
        %1418 = arith.addi %arg165, %arg162 : index
        memref.store %cst_3, %alloc_1710[%c0, %1416, %1417, %1418] : memref<1x512x7x7xf32>
    }
    scf.reduce
}
scf.reduce
}
```

Figure 4: Optimized kernel generation showing improved block size mapping.

The figure above demonstrates the result of this modification. The block size is now mapped to {1, 32, 8}, providing significantly higher occupancy compared to the baseline.

4 Switch to IREE

4.1 Transition to IREE

We transitioned our experimental framework to IREE to remove issues faced when executing binaries generated directly through the standard MLIR pipeline. While IREE offered a more streamlined workflow for binary generation and runtime verification, the migration showed a fundamental mismatch between the upstream MLIR and IREE code generation pipelines:

- **Incompatibility with Upstream Mechanisms:** We initially targeted the upstream MLIR Parallel LoopMapper (implemented in `ParallelLoopMapper.cpp`). This pass typically annotates `scf.parallel` operations with mapping attributes, allowing the `convert-parallel-loops-to-gpu` pass to lower them to `gpu.launch`.
- **IREE Codegen Architecture:** However, we discovered that IREE’s default GPU pipelines do not rely on the `scf.parallel` lowering path. Instead, IREE primarily utilizes `scf.forall` to handle workgroup and thread semantics, employing its own custom passes for GPU distribution and tiling.
- **Adaptation of Strategy:** Because IREE effectively bypassed the upstream mapper we had modified, we could not use it to benchmark our optimization. Consequently, we shifted our focus to modifying IREE’s internal LLVMGPU kernel configuration to optimize tile sizes and thread mapping directly within their custom infrastructure.

4.2 Analysis

First we profiled the vanilla IREE based resnet50 executable using ncu and obtained the following results (shown in the figure):

We also ran `iree-benchmark-module` on the module to get the actual run time without the overhead of ncu, and got the following time: 19.7ms

We could make the following inferences:

- Compute throughput was very low for most kernels.
- The slowest kernels which contributed the most to total time are the matrix multiplication kernels.

ID	Estimated Sp Function Name	D _e	Duration [ms] (31.50 ms)	Runtime Improvement [ms] (24.14 ms)	Compute Throughput [%]	Memory Throughput [%]	# Registers [register/thread]	Grid Size	Block Size [block]
0	62.59 forward_dispatch_0_slow_memcpy	...	0.01	0.00	12.78	35.77	16	2688,	1, -
1	35.01 forward_dispatch_1_conv_64x12x112x3x7x7_f32	...	0.17	0.06	97.12	97.12	58	25688,	1, -
2	35.11 forward_dispatch_2_conv_64x5x5x3x3_f32	...	0.01	0.00	12.23	77.53	20	6272,	1, -
3	86.71 forward_dispatch_3_matmul_like_64x6x6x6x4_f32	...	0.33	0.28	4.41	53.51	63	128,	1, -
4	48.15 forward_dispatch_4_conv_64x56x56x4x3x3_f32	...	0.23	0.11	98.17	95.29	34	6272,	1, -
5	48.66 forward_dispatch_5_matmul_like_256x3136x64_f32	...	0.25	0.12	19.62	66.02	55	208,	1, -
6	67.29 forward_dispatch_6_matmul_like_256x3136x64_f32	...	0.15	0.10	34.01	69.46	55	200,	1, -
7	86.43 forward_dispatch_7_matmul_like_64x56x56x256_f32	...	1.39	1.20	3.99	50.57	63	128,	1, -
8	47.75 forward_dispatch_8_conv_64x56x56x4x3x3_f32	...	0.23	0.11	89.46	94.56	34	6272,	1, -
9	67.09 forward_dispatch_9_matmul_like_256x3136x64_f32	...	0.15	0.10	53.97	69.25	55	208,	1, -
10	86.28 forward_dispatch_10_matmul_like_4x56x56x256_f32	...	1.41	1.22	3.94	50.12	63	128,	1, -
11	48.03 forward_dispatch_11_conv_64x56x56x4x3x3_f32	...	0.22	0.11	89.94	95.05	34	6272,	1, -
12	67.39 forward_dispatch_12_matmul_like_256x3136x64_f32	...	0.15	0.10	34.13	69.56	55	208,	1, -
13	84.66 forward_dispatch_13_matmul_like_128x56x56x256_f32	...	3.26	2.75	3.41	52.64	63	256,	1, -
14	59.86 forward_dispatch_14_conv_128x28x28x128x3x3_f32	...	0.26	0.16	64.84	96.34	36	3136,	1, -
15	82.46 forward_dispatch_15_matmul_like_512x28x28x56_f32	...	2.04	1.68	6.20	61.31	54	512,	1, -
16	68.18 forward_dispatch_16_matmul_like_512x784x128_f32	...	0.16	0.10	31.25	62.12	55	112,	1, -
17	73.81 forward_dispatch_17_matmul_like_128x28x28x512_f32	...	0.33	0.24	19.05	40.63	63	128,	1, -
18	64.14 forward_dispatch_18_conv_128x28x28x128x3_f32	...	0.26	0.17	77.69	92.55	36	3136,	1, -
19	59.75 forward_dispatch_19_matmul_like_512x784x128_f32	...	0.16	0.10	31.00	61.68	65	112,	1, -
20	83.30 forward_dispatch_17_matmul_like_128x28x28x512_f32	...	0.32	0.26	19.93	42.29	63	128,	1, -
21	64.06 forward_dispatch_18_conv_128x28x28x128x3_f32	...	0.26	0.17	77.12	92.40	36	3136,	1, -
22	68.04 forward_dispatch_19_matmul_like_512x784x128_f32	...	0.16	0.10	31.15	61.98	55	112,	1, -
23	82.08 forward_dispatch_17_matmul_like_128x28x28x512_f32	...	0.33	0.27	19.44	40.13	63	128,	1, -
24	63.73 forward_dispatch_18_conv_128x28x28x128x3_f32	...	0.26	0.17	76.72	91.94	36	3136,	1, -
25	68.14 forward_dispatch_19_matmul_like_512x784x128_f32	...	0.16	0.10	31.21	62.08	65	112,	1, -
26	77.32 forward_dispatch_26_matmul_like_256x28x28x512_f32	...	1.54	1.19	8.19	21.92	63	256,	1, -
27	74.84 forward_dispatch_27_conv_256x14x14x256x3_f32	...	0.53	0.40	31.55	88.60	36	1568,	1, -
28	67.46 forward_dispatch_28_matmul_like_1024x14x14x512_f32	...	0.87	0.59	23.97	63.63	40	1024,	1, -
29	56.38 forward_dispatch_29_matmul_like_1024x196x256_f32	...	0.28	0.11	26.44	43.62	55	64,	1, -
30	89.73 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	0.64	0.58	8.25	10.27	44	256,	1, -
31	74.75 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.52	0.39	38.70	87.81	34	1568,	1, -
32	56.48 forward_dispatch_32_matmul_like_1024x196x256_f32	...	0.28	0.11	26.42	43.60	55	64,	1, -
33	89.65 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	0.64	0.57	8.35	10.35	44	256,	1, -
44	74.49 forward_dispatch_31_conv_256x14x14x256x3_f32	A-F-A	0.54	0.44	38.93	88.30	34	1568,	1, -
23	82.08 forward_dispatch_17_matmul_like_128x28x28x512_f32	...	0.33	0.27	19.44	40.13	63	128,	1, -
24	63.73 forward_dispatch_18_conv_128x28x28x128x3_f32	...	0.26	0.17	76.72	91.94	36	3136,	1, -
25	68.14 forward_dispatch_19_matmul_like_512x784x128_f32	...	0.16	0.10	31.21	62.08	65	112,	1, -
26	77.32 forward_dispatch_26_matmul_like_256x28x28x512_f32	...	1.54	1.19	8.19	21.92	63	256,	1, -
27	74.84 forward_dispatch_27_conv_256x14x14x256x3_f32	...	0.53	0.40	31.55	88.60	36	1568,	1, -
28	67.46 forward_dispatch_28_matmul_like_1024x14x14x512_f32	...	0.87	0.59	23.97	63.63	40	1024,	1, -
29	56.38 forward_dispatch_29_matmul_like_1024x196x256_f32	...	0.28	0.11	26.44	43.62	55	64,	1, -
30	89.73 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	0.64	0.58	8.25	10.27	44	256,	1, -
31	74.75 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.52	0.39	38.70	87.81	34	1568,	1, -
32	56.48 forward_dispatch_32_matmul_like_1024x196x256_f32	...	0.28	0.11	26.42	43.60	55	64,	1, -
33	89.65 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	0.64	0.57	8.35	10.35	44	256,	1, -
34	74.31 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.52	0.39	38.93	88.30	34	1568,	1, -
35	56.34 forward_dispatch_32_matmul_like_1024x14x14x512_f32	...	0.28	0.11	26.46	43.66	55	64,	1, -
36	89.84 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.65	0.58	8.18	10.16	44	256,	1, -
37	75.80 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.51	0.39	39.25	89.04	34	1568,	1, -
38	56.40 forward_dispatch_32_matmul_like_1024x196x256_f32	...	0.28	0.11	26.41	43.60	55	64,	1, -
39	89.85 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	0.65	0.58	8.19	10.15	44	256,	1, -
40	74.83 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.52	0.39	38.74	87.90	34	1568,	1, -
41	58.37 forward_dispatch_32_matmul_like_1024x196x256_f32	...	0.28	0.11	26.45	43.63	55	64,	1, -
42	89.81 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	0.65	0.58	8.19	10.19	44	256,	1, -
43	74.95 forward_dispatch_31_conv_256x14x14x256x3_f32	...	0.52	0.39	38.80	88.04	34	1568,	1, -
44	56.45 forward_dispatch_32_matmul_like_1024x196x256_f32	...	0.28	0.11	26.39	43.55	55	64,	1, -
45	77.80 forward_dispatch_45_matmul_like_512x14x14x1024_f32	...	0.58	0.45	18.39	22.20	46	512,	1, -
46	83.68 forward_dispatch_46_conv_512x7x7x512x3x3_f32	...	1.74	1.46	9.61	95.09	34	784,	1, -
47	65.43 forward_dispatch_47_matmul_like_2048x49x512_f32	...	0.88	0.58	48.18	48.18	48	2048,	1, -
48	74.87 forward_dispatch_48_matmul_like_2048x49x512_f32	...	0.26	0.19	24.56	25.13	55	64,	1, -
49	62.45 forward_dispatch_49_matmul_like_512x7x7x512x3_f32	...	0.56	0.35	37.55	37.55	48	512,	1, -
50	81.84 forward_dispatch_50_conv_512x7x7x512x3x3_f32	...	1.78	1.45	11.30	93.16	36	784,	1, -
51	74.30 forward_dispatch_51_matmul_like_2048x49x512_f32	...	0.25	0.19	25.12	25.70	55	64,	1, -
52	62.58 forward_dispatch_49_matmul_like_512x7x7x512x3_f32	...	0.57	0.35	37.42	37.42	48	512,	1, -
53	81.51 forward_dispatch_50_conv_512x7x7x512x3x3_f32	...	1.78	1.45	11.28	92.97	36	784,	1, -
54	69.22 forward_dispatch_54_matmul_like_2048x49x512_f32	...	0.29	0.20	21.80	30.78	54	64,	1, -
55	79.87 forward_dispatch_55_reduction_2048x49_f32	...	0.01	0.01	1.44	20.13	77	16,	1, -
56	90.00 forward_dispatch_56_matvec_like_1000x2048_f32	...	0.03	0.01	7.57	87.81	48	1000,	1, -

Figure 5: Default iree performance with O3

- The thread configuration in most of the blocks is (32, 8, 1) and some of them (mostly convolution kernels) are (32, 1, 1)
- Most kernels have a very high percentage of uncoalesced accesses.

4.3 Trying to improve throughput of matrix multiplication kernels

First, we focused on trying to improve the throughput of matrix multiplication kernels as they were the bottleneck. We looked through the code in the "KernelConfig.cpp" file, and observed the following:

- There is a fixed set of tile sizes and thread block mapping in the following function:

```
static SmallVector<TileWorkgroupSizePair>
getMatmulConfig(IREE::GPU::TargetAttr target) {
    SmallVector<TileWorkgroupSizePair> tileSizes;
    // Pick tile size so that M*K and K*N divisible by wgSize * \*vecSize*\4.
    // This way workgroup memory copy don't need to be masked. Once we support
    // masked load we can get performance out of more configuration.

    // Make use of the full subgroup when possible.
    if (target.getPreferredSubgroupSize() == 64) {
        tileSizes.push_back(TileWorkgroupSizePair({{64, 128, 64}, {64, 16, 1}, 1}));
    }

    llvm::append_values(tileSizes,
                        TileWorkgroupSizePair({{32, 128, 32}, {32, 8, 1}, 1}),
                        TileWorkgroupSizePair({{128, 64, 8}, {16, 8, 1}, 1}),
                        TileWorkgroupSizePair({{16, 256, 32}, {64, 2, 1}, 1}),
                        TileWorkgroupSizePair({{8, 32, 32}, {8, 8, 1}, 1}),
                        TileWorkgroupSizePair({{32, 128, 4}, {32, 8, 1}, 1}),
                        TileWorkgroupSizePair({{8, 128, 4}, {32, 1, 1}, 1}),
                        TileWorkgroupSizePair({{16, 64, 4}, {16, 2, 1}, 1}),
                        TileWorkgroupSizePair({{1, 128, 8}, {32, 1, 1}, 1}));

    return tileSizes;
}
```

Here, the first three values in TileWorkgroupSizePair correspond to the M, N and K dimensions of the tiles, while the next three correspond to the x, y and z block dimensions.

- Also, in the function which assigns the tiling dimensions to each kernel (setContractConfig), we noticed that if the boolean "isStaticSize" was true, which was not the case for all our kernels, it would simply use the first element of the tileSizes variable to assign tile size. This would be:

```
TileWorkgroupSizePair({{32, 128, 32}, {32, 8, 1}, 1})  
which is not really optimal for resnet50.
```

- Our intuition was that a square tile size with width 16 will probably perform the best on our given GPU.

- So, we added this tile size to the top of the list:

```
TileWorkgroupSizePair({{16, 16, 16}, {16, 16, 1}, 1})
```

- This gives a perfectly square 16x16 tile and also assigns a 16x16 thread block to each tile.

4.3.1 Results

With this simple optimization, The runtime of the module that we got with iree-benchmark-module went down to 10.3ms. The ncu report of that run is shown in the figures below:

4.3.2 Observations

- We observed that all the matrix multiplication kernels were quite fast now, with much higher memory and compute throughputs overall.
- However, now the bottleneck was the convolution kernels.
- So, next we focused on finding out why the convolution kernels were so slow, with very low compute throughput and a lot of uncoalesced accesses.

4.4 Improving the convolution kernels

With a lot of debugging and digging through the code, we found that the tile sizes for the convolutions were being mapped by the "setTileAndFuseLoweringConfig" function in ConfigUtils.cpp. The part of this function that actually does this mapping is the "distributeToThreads" lambda in this function. Here, we can take an example of a particular convolution kernel to explain what this does:

- We'll look at the `forward_dispatch_46_conv_512x7x7x512x3x3_f32` kernel, which outputs a tensor with dimension 512x7x7, with 512x3x3 kernels for each input channel.
- First, this lambda finds the "partitionable" loops in this kernel, which in this case are the ones with dimensions 512, 7 and 7.
- Then, it goes from the innermost dimension to the outermost dimension and tries to assign the highest number of threads (powers of 2) to each dimension.
- However, there is also a "lossFactor" parameter in this lambda, and it will only assign threads to a dimension if the fraction of idle threads are lesser than or equal to $\frac{1}{\text{lossFactor}}$
- However, if we look at where this lambda is called, we see that it first tries to assign threads with a 0 loss factor.
- This will actually result in threads being tiled only across the 512 dimension, since the dimension of 7 will always result in some threads being idle.
- Instead, we thought that maybe distributing threads across the 7x7 dimensions, it would give better throughput.
- So, we changed the first call to this lambda to take a lossFactor of 4, so that it allows $\frac{1}{4}$ of the threads to be idle. The results for that are shown below:

4.4.1 Results

From iree-benchmark-module, we got a runtime of 6.84ms.

4.4.2 Observations

- Here, we observe that the compute throughput and coalescing has significantly improved for the convolution kernels.
- We tried to increase the number of threads per block for the convolution kernels, but they were giving worse results.

ID	Estimated Sp Function Name	De	Duration [ms] (15.95 ms)	Runtime Improvement [ms] (9.79 ms)	Compute Throughput [%]	Memory Throughput [%]	# Registers [register/thread]	Grid Size	Block Size [block]
0	48_30 forward_dispatch_0_slow_memcpy	-	0.02	0.01	51.70	26.90	16	9480, 1, -	16, 16, -
1	34_97 forward_dispatch_1_conv_64x12x12x3x7x7_f32	-	0.17	0.06	97.00	97.00	58	25688, 1, -	32, 1, -
2	34_58 forward_dispatch_2_conv_64x56x56x3x3_f32	-	0.01	0.00	12.16	77.53	20	6272, 1, -	32, 1, -
3	25_00 forward_dispatch_3_matmul_like_64x6x6x6x6x4_f32	-	0.04	0.01	39.71	88.65	48	1024, 1, -	16, 16, -
4	47_67 forward_dispatch_4_conv_64x56x6x4x3x3_f32	-	0.23	0.11	89.29	94.34	34	6272, 1, -	32, 1, -
5	9_42 forward_dispatch_5_matmul_like_256x3136x64_f32	-	0.06	0.01	94.16	94.16	48	3136, 1, -	16, 16, -
6	11_65 forward_dispatch_6_matmul_like_256x3136x64_f32	-	0.07	0.01	93.75	93.75	41	3136, 1, -	16, 16, -
7	27_47 forward_dispatch_7_matmul_like_64x56x56x256_f32	-	0.15	0.04	38.94	85.62	40	1024, 1, -	16, 16, -
8	47_39 forward_dispatch_8_conv_64x56x56x4x3x3_f32	-	0.23	0.11	88.81	93.78	34	6272, 1, -	32, 1, -
9	18_65 forward_dispatch_9_matmul_like_256x3136x64_f32	-	0.07	0.01	94.14	94.14	48	3136, 1, -	16, 16, -
10	27_13 forward_dispatch_10_matmul_like_64x56x56x256_f32	-	0.15	0.04	39.19	86.34	48	1024, 1, -	16, 16, -
11	47_63 forward_dispatch_11_conv_64x56x56x64x3x3_f32	-	0.23	0.11	89.23	94.26	34	6272, 1, -	32, 1, -
12	10_66 forward_dispatch_9_matmul_like_256x3136x64_f32	-	0.07	0.01	94.27	94.27	48	3136, 1, -	16, 16, -
13	30_53 forward_dispatch_13_matmul_like_128x56x56x256_f32	-	0.30	0.09	37.77	91.01	48	2048, 1, -	16, 16, -
14	58_96 forward_dispatch_14_conv_28x28x128x3x3_f32	-	0.26	0.16	63.85	88.99	36	3136, 1, -	32, 1, -
15	48_38 forward_dispatch_15_matmul_like_512x28x28x32_f32	-	0.66	0.32	17.56	62.54	40	2048, 1, -	16, 16, -
16	18_49 forward_dispatch_16_matmul_like_512x78x128_f32	-	0.06	0.01	93.21	93.21	48	1568, 1, -	16, 16, -
17	58_88 forward_dispatch_17_matmul_like_128x28x28x12_f32	-	0.11	0.06	58.61	61.16	48	512, 1, -	16, 16, -
18	64_88 forward_dispatch_18_conv_28x28x128x3x3_f32	-	0.26	0.17	77.89	92.32	36	3136, 1, -	32, 1, -
19	9_95 forward_dispatch_19_matmul_like_512x78x128_f32	-	0.06	0.01	93.25	93.25	44	1568, 1, -	16, 16, -
20	50_00 forward_dispatch_17_matmul_like_128x28x28x12_f32	-	0.12	0.06	48.45	58.56	40	512, 1, -	16, 16, -
21	63_36 forward_dispatch_18_conv_28x28x128x3x3_f32	-	0.26	0.17	76.20	91.39	36	3136, 1, -	32, 1, -
22	9_97 forward_dispatch_19_matmul_like_512x78x128_f32	-	0.06	0.01	93.42	93.42	44	1568, 1, -	16, 16, -
23	58_00 forward_dispatch_17_matmul_like_128x28x28x12_f32	-	0.12	0.06	47.93	58.83	40	512, 1, -	16, 16, -
24	63_87 forward_dispatch_18_conv_18_conv_28x28x128x3x3_f32	-	0.26	0.17	76.82	92.19	36	3136, 1, -	32, 1, -
25	9_95 forward_dispatch_19_matmul_like_512x78x128_f32	-	0.06	0.01	93.26	93.26	44	1568, 1, -	16, 16, -
26	36_76 forward_dispatch_26_matmul_like_256x28x28x12_f32	-	0.23	0.08	50.71	63.24	48	1024, 1, -	16, 16, -
27	74_98 forward_dispatch_27_conv_256x14x14x256x3x3_f32	-	0.53	0.40	51.63	88.77	36	1568, 1, -	32, 1, -
28	55_31 forward_dispatch_28_matmul_like_1024x14x14x512_f32	-	0.70	0.39	18.24	68.67	39	1024, 1, -	16, 16, -
29	23_01 forward_dispatch_29_matmul_like_1024x196x256_f32	-	0.07	0.02	88.81	88.81	48	832, 1, -	16, 16, -
30	26_38 forward_dispatch_30_matmul_like_256x14x14x1024_f32	-	0.09	0.02	73.62	73.62	38	256, 1, -	16, 16, -
31	75_75 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.51	0.39	39.21	88.98	34	1568, 1, -	32, 1, -
32	22_65 forward_dispatch_32_matmul_like_1024x196x256_f32	-	0.07	0.01	88.95	88.95	42	832, 1, -	16, 16, -
33	23_82 forward_dispatch_30_matmul_like_256x14x14x1024_f32	-	0.08	0.02	76.18	76.18	38	256, 1, -	16, 16, -
44	44_44 forward_dispatch_44_conv_256x14x14x256x3x3_f32	A_Pt	0.44	0.44	44.44	44.44	44	44, 1, -	44, 1, -
23	58_00 forward_dispatch_17_matmul_like_128x28x28x12_f32	-	0.12	0.06	47.93	58.83	40	512, 1, -	16, 16, -
24	63_87 forward_dispatch_18_conv_128x28x28x128x3x3_f32	-	0.26	0.17	76.82	92.19	36	3136, 1, -	32, 1, -
25	9_95 forward_dispatch_19_matmul_like_512x78x128_f32	-	0.06	0.01	93.26	93.26	44	1568, 1, -	16, 16, -
26	36_76 forward_dispatch_26_matmul_like_256x28x28x12_f32	-	0.23	0.08	50.71	63.24	48	1024, 1, -	16, 16, -
27	74_98 forward_dispatch_27_conv_256x14x14x256x3x3_f32	-	0.53	0.40	51.63	88.77	36	1568, 1, -	32, 1, -
28	55_31 forward_dispatch_28_matmul_like_1024x14x14x512_f32	-	0.70	0.39	18.24	68.67	39	1024, 1, -	16, 16, -
29	23_01 forward_dispatch_29_matmul_like_1024x196x256_f32	-	0.07	0.02	88.81	88.81	48	832, 1, -	16, 16, -
30	26_38 forward_dispatch_30_matmul_like_256x14x14x1024_f32	-	0.09	0.02	73.62	73.62	38	256, 1, -	16, 16, -
31	75_75 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.51	0.39	39.21	88.98	34	1568, 1, -	32, 1, -
32	22_65 forward_dispatch_32_matmul_like_1024x196x256_f32	-	0.07	0.01	88.95	88.95	42	832, 1, -	16, 16, -
33	23_82 forward_dispatch_30_matmul_like_256x14x14x1024_f32	-	0.08	0.02	76.18	76.18	38	256, 1, -	16, 16, -
34	75_79 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.51	0.39	39.25	88.94	34	1568, 1, -	32, 1, -
35	22_71 forward_dispatch_32_matmul_like_1024x196x256_f32	-	0.07	0.01	88.27	88.27	42	832, 1, -	16, 16, -
36	25_69 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.09	0.02	74.31	74.31	38	256, 1, -	16, 16, -
37	75_67 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.51	0.39	39.18	88.90	34	1568, 1, -	32, 1, -
38	22_88 forward_dispatch_32_matmul_like_1024x196x256_f32	-	0.06	0.01	88.92	88.92	42	832, 1, -	16, 16, -
39	26_07 forward_dispatch_30_matmul_like_256x14x14x1024_f32	-	0.09	0.02	73.93	73.93	38	256, 1, -	16, 16, -
40	75_73 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.51	0.39	39.21	88.97	34	1568, 1, -	32, 1, -
41	22_57 forward_dispatch_32_matmul_like_1024x196x256_f32	-	0.07	0.01	87.71	87.71	42	832, 1, -	16, 16, -
42	24_99 forward_dispatch_30_matmul_like_256x14x14x1024_f32	-	0.09	0.02	75.01	75.01	38	256, 1, -	16, 16, -
43	74_92 forward_dispatch_31_conv_256x14x14x256x3x3_f32	-	0.52	0.39	38.79	88.81	34	1568, 1, -	32, 1, -
44	22_57 forward_dispatch_32_matmul_like_1024x196x256_f32	-	0.07	0.01	87.72	87.72	42	832, 1, -	16, 16, -
45	50_80 forward_dispatch_45_matmul_like_512x14x14x1024_f32	-	0.16	0.08	88.59	88.59	38	512, 1, -	16, 16, -
46	83_73 forward_dispatch_46_conv_7512x7x512x3x3_f32	-	1.74	1.46	9.62	95.17	34	784, 1, -	32, 1, -
47	57_95 forward_dispatch_47_matmul_like_2048x97x1024_f32	-	0.30	0.17	48.75	79.36	39	2048, 1, -	16, 16, -
48	50_80 forward_dispatch_48_matmul_like_2048x49512_f32	-	0.07	0.04	79.17	79.17	40	512, 1, -	16, 16, -
49	61_84 forward_dispatch_49_matmul_like_512x7x2048_f32	-	0.19	0.12	38.16	38.16	38	512, 1, -	16, 16, -
50	82_05 forward_dispatch_50_conv_512x7x7x12x3x3_f32	-	1.77	1.46	11.32	93.37	36	784, 1, -	32, 1, -
51	50_80 forward_dispatch_51_matmul_like_2048x49512_f32	-	0.07	0.04	79.86	79.86	40	512, 1, -	16, 16, -
52	62_12 forward_dispatch_49_matmul_like_512x7x7x2048_f32	-	0.19	0.12	37.88	37.88	38	512, 1, -	16, 16, -
53	81_49 forward_dispatch_50_conv_512x7x7x12x3x3_f32	-	1.78	1.45	11.26	92.87	36	784, 1, -	32, 1, -
54	50_80 forward_dispatch_54_matmul_like_2048x49512_f32	-	0.07	0.04	77.36	77.36	40	512, 1, -	16, 16, -
55	88_48 forward_dispatch_55_reduction_2048x49_f32	-	0.01	0.01	1.39	19.52	77	16, 1, -	32, 1, -
56	88_00 forward_dispatch_56_mvavve_like_1000x2048_f32	-	0.01	0.01	7.49	87.51	48	1000, 1, -	128, 1, -

Figure 6: Report for run with mat-mult tilesize optimization

ID	Estimated Sp Function Name	D ₀	Duration [us] (9,529.47 us)	Runtime Improvement [us] (3,641.18 us)	Compute Throughput [%]	Memory Throughput [%]	# Registers	register/thread	Grid Size	Block Size [block]
0	48.28 forward_dispatch_0_slow_memcpy	...	23.65	11.42	51.72	26.98	16	9400	1, -	16, 16, -
1	34.96 forward_dispatch_1_conv_64x12x12x3x7x7_f32	...	170.66	59.67	96.98	96.98	58	25088	1, -	32, 1, -
2	25.10 forward_dispatch_2_conv_64x56x56x3x3_f32	...	12.51	3.14	12.96	77.53	24	7168	1, -	32, 1, -
3	25.00 forward_dispatch_3_matmul_like_64x6x6x6x6x4_f32	...	39.01	9.75	39.64	81.17	40	1024	1, -	16, 16, -
4	30.44 forward_dispatch_4_conv_64x56x6x6x4x3_f32	...	243.58	74.16	94.83	94.83	38	7168	1, -	32, 1, -
5	9.44 forward_dispatch_5_matmul_like_256x3136x64_f32	...	61.57	5.81	94.36	94.36	40	3136	1, -	16, 16, -
6	11.64 forward_dispatch_6_matmul_like_256x3136x64_f32	...	68.03	7.92	93.66	93.66	41	3136	1, -	16, 16, -
7	26.27 forward_dispatch_7_matmul_like_64x56x56x256_f32	...	148.19	38.93	38.68	84.14	40	1024	1, -	16, 16, -
8	30.53 forward_dispatch_8_conv_64x56x56x4x3_f32	...	242.85	74.14	95.09	95.09	38	7168	1, -	32, 1, -
9	18.66 forward_dispatch_9_matmul_like_256x3136x64_f32	...	65.41	6.97	94.22	94.22	40	3136	1, -	16, 16, -
10	27.31 forward_dispatch_10_matmul_like_64x56x56x256_f32	...	147.81	48.37	38.76	85.62	40	1024	1, -	16, 16, -
11	30.54 forward_dispatch_11_conv_64x56x56x64x3_f32	...	242.72	74.13	95.13	95.13	38	7168	1, -	32, 1, -
12	10.65 forward_dispatch_9_matmul_like_256x3136x64_f32	...	65.47	6.97	94.14	94.14	40	3136	1, -	16, 16, -
13	30.29 forward_dispatch_13_matmul_like_128x56x56x256_f32	...	306.43	92.81	37.35	89.07	40	2048	1, -	16, 16, -
14	33.24 forward_dispatch_14_conv_28x28x128x3x3_f32	...	238.46	79.26	88.59	88.59	40	3584	1, -	32, 1, -
15	48.46 forward_dispatch_15_matmul_like_512x28x28x32_f32	...	647.17	313.59	17.80	62.87	40	2048	1, -	16, 16, -
16	18.50 forward_dispatch_16_matmul_like_512x784x128_f32	...	64.16	6.74	93.29	93.29	40	1568	1, -	16, 16, -
17	58.86 forward_dispatch_17_matmul_like_128x28x28x12_f32	...	111.38	55.65	51.94	62.82	40	512	1, -	16, 16, -
18	29.45 forward_dispatch_18_conv_128x28x128x3x3_f32	...	250.82	73.87	91.84	91.84	37	3584	1, -	32, 1, -
19	9.98 forward_dispatch_19_matmul_like_512x784x128_f32	...	63.33	6.32	93.50	93.50	44	1568	1, -	16, 16, -
20	50.00 forward_dispatch_17_matmul_like_128x28x28x12_f32	...	119.14	59.57	48.66	58.83	40	512	1, -	16, 16, -
21	29.63 forward_dispatch_18_conv_128x28x128x3x3_f32	...	249.22	73.85	92.41	92.41	37	3584	1, -	32, 1, -
22	9.97 forward_dispatch_19_matmul_like_512x784x128_f32	...	63.07	6.29	93.49	93.49	44	1568	1, -	16, 16, -
23	50.00 forward_dispatch_17_matmul_like_128x28x28x12_f32	...	116.19	58.10	49.80	60.21	40	512	1, -	16, 16, -
24	29.84 forward_dispatch_18_conv_128x28x128x3x3_f32	...	248.03	74.00	93.04	93.04	37	3584	1, -	32, 1, -
25	9.99 forward_dispatch_19_matmul_like_512x784x128_f32	...	62.91	6.29	93.67	93.67	44	1568	1, -	16, 16, -
26	35.09 forward_dispatch_26_matmul_like_256x28x28x12_f32	...	232.51	81.58	49.63	64.91	40	1024	1, -	16, 16, -
27	37.62 forward_dispatch_27_conv_256x14x12x56x3x3_f32	...	232.35	87.41	82.43	82.43	40	1792	1, -	32, 1, -
28	55.57 forward_dispatch_28_matmul_like_1024x14x14x512_f32	...	696.58	387.12	18.40	76.99	39	1024	1, -	16, 16, -
29	22.99 forward_dispatch_29_matmul_like_1024x196x256_f32	...	65.34	15.02	88.72	88.72	40	832	1, -	16, 16, -
30	25.29 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	66.08	21.77	74.71	74.71	38	256	1, -	16, 16, -
31	50.00 forward_dispatch_31_conv_256x14x14x256x3x3_f32	...	276.22	138.11	83.22	83.22	38	1792	1, -	32, 1, -
32	22.77 forward_dispatch_32_matmul_like_1024x196x256_f32	...	64.93	14.79	88.52	88.52	42	832	1, -	16, 16, -
33	25.57 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	66.37	22.08	74.43	74.43	38	256	1, -	16, 16, -
44	44.44 forward_dispatch_44_conv_128x128x128x3x3_f32	...	75.44	44.44	44.44	44.44	44	44	44, -	44, 4
23	50.00 forward_dispatch_17_matmul_like_128x28x28x12_f32	...	116.19	58.10	49.80	60.21	40	512	1, -	16, 16, -
24	29.84 forward_dispatch_18_conv_128x28x128x3x3_f32	...	248.03	74.00	93.04	93.04	37	3584	1, -	32, 1, -
25	9.99 forward_dispatch_19_matmul_like_512x784x128_f32	...	62.91	6.29	93.67	93.67	44	1568	1, -	16, 16, -
26	35.09 forward_dispatch_26_matmul_like_256x28x28x12_f32	...	232.51	81.58	49.63	64.91	40	1024	1, -	16, 16, -
27	37.62 forward_dispatch_27_conv_256x14x12x56x3x3_f32	...	232.35	87.41	82.43	82.43	40	1792	1, -	32, 1, -
28	55.57 forward_dispatch_28_matmul_like_1024x14x14x512_f32	...	696.58	387.12	18.40	76.99	39	1024	1, -	16, 16, -
29	22.99 forward_dispatch_29_matmul_like_1024x196x256_f32	...	65.34	15.02	88.72	88.72	40	832	1, -	16, 16, -
30	25.29 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	66.08	21.77	74.71	74.71	38	256	1, -	16, 16, -
31	50.00 forward_dispatch_31_conv_256x14x14x256x3x3_f32	...	276.22	138.11	83.22	83.22	38	1792	1, -	32, 1, -
32	22.77 forward_dispatch_32_matmul_like_1024x196x256_f32	...	64.93	14.79	88.52	88.52	42	832	1, -	16, 16, -
33	25.57 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	66.37	22.08	74.43	74.43	38	256	1, -	16, 16, -
34	22.77 forward_dispatch_32_matmul_like_1024x196x256_f32	...	277.92	63.28	82.73	82.73	38	1792	1, -	32, 1, -
35	22.63 forward_dispatch_32_matmul_like_1024x196x256_f32	...	65.38	14.80	87.97	87.97	42	832	1, -	16, 16, -
36	23.73 forward_dispatch_32_matmul_like_1024x196x256_f32	...	84.26	19.99	76.27	76.27	38	256	1, -	16, 16, -
37	50.00 forward_dispatch_31_conv_256x14x14x256x3x3_f32	...	278.37	139.18	82.58	82.58	38	1792	1, -	32, 1, -
38	22.86 forward_dispatch_32_matmul_like_256x14x14x256x3x3_f32	...	64.61	14.77	88.86	88.86	42	832	1, -	16, 16, -
39	24.37 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	85.02	20.72	75.63	75.63	38	256	1, -	16, 16, -
40	50.00 forward_dispatch_31_conv_256x14x14x256x3x3_f32	...	277.76	138.88	82.72	82.72	38	1792	1, -	32, 1, -
41	22.94 forward_dispatch_32_matmul_like_1024x196x256_f32	...	64.26	14.74	89.17	89.17	42	832	1, -	16, 16, -
42	24.69 forward_dispatch_30_matmul_like_256x14x14x1024_f32	...	85.22	21.04	75.31	75.31	38	256	1, -	16, 16, -
43	50.00 forward_dispatch_31_conv_256x14x14x256x3x3_f32	...	280.86	140.43	81.80	81.80	38	1792	1, -	32, 1, -
44	22.84 forward_dispatch_32_matmul_like_1024x196x256_f32	...	64.88	14.80	88.79	88.79	42	832	1, -	16, 16, -
45	50.00 forward_dispatch_45_matmul_like_512x14x14x1024_f32	...	158.50	79.25	81.11	81.11	38	512	1, -	16, 16, -
46	32.48 forward_dispatch_46_conv_512x7x7x512x3x3_f32	...	324	105.23	67.50	74.22	42	1024	1, -	32, 1, -
47	57.90 forward_dispatch_47_matmul_like_2048x9x7x1024_f32	...	297.09	172.80	49.43	86.38	39	2048	1, -	16, 16, -
48	50.00 forward_dispatch_48_matmul_like_2048x49x512_f32	...	77.92	38.96	74.49	74.49	40	512	1, -	16, 16, -
49	61.63 forward_dispatch_49_matmul_like_512x7x7x2048_f32	...	199.78	117.58	38.37	38.37	38	512	1, -	16, 16, -
50	34.13 forward_dispatch_50_conv_512x7x7x512x3x3_f32	...	328.51	112.12	79.92	79.92	44	1024	1, -	32, 1, -
51	50.00 forward_dispatch_51_matmul_like_2048x49x512_f32	...	72.70	36.35	79.77	79.77	40	512	1, -	16, 16, -
52	62.11 forward_dispatch_49_matmul_like_512x7x7x2048_f32	...	193.22	120.01	37.89	37.89	38	512	1, -	16, 16, -
53	33.58 forward_dispatch_50_conv_512x7x7x512x3x3_f32	...	333.82	112.08	78.63	78.63	44	1024	1, -	32, 1, -
54	50.00 forward_dispatch_54_matmul_like_2048x49x512_f32	...	77.63	38.82	74.28	74.28	40	512	1, -	16, 16, -
55	66.85 forward_dispatch_55_reduction_2048x49_f32	...	7.33	4.81	2.56	33.95	38	64	1, -	32, 1, -
56	50.00 forward_dispatch_56_mavvee_like_1000x2048_f32	...	27.46	15.73	7.51	87.79	40	1000	1, -	128, 1, -

Figure 7: Report for run with both mat-mul and conv tilesize optimizations

Configuration	Runtime (ms)	Speedup
Original IREE compile	19.7	-
With 16×16 matmul	10.3	1.91x
With optimized convolutions	6.84	2.88x

Table 1: Final results

5 Final results

Overall, we see a 2.88x improvement over the default iree compiler with our improvements. The kernels went from mostly having low compute and memory throughput to now most of them having close to maximum compute and memory throughput. We also see that the number of uncoalesced accesses decreased significantly.

6 Challenges Faced

- **MLIR Documentation Issues:** Navigating through the MLIR passes was difficult because the documentation was often incomplete or unclear.
- **IREE Complexity:** Finding the right place to make modifications in IREE was even harder, as there was almost no documentation for its internal passes.
- **Execution Delays:** We faced significant issues when trying to run the generated executables. It took approximately two weeks of debugging to resolve these errors and successfully run the code.

7 Future Work

- Develop an interface that allows users to manually configure the tile sizes and thread block mappings.
- Extend the current implementation to support a wider variety of kernel sizes and input shapes.